

Benchmarking

CS553 Assignment #1

Maulik Patel A20334253

The assignment is about benchmarking different parts of computer. I have benchmarked three parts of the computer as follows

1. Network Performance.
2. Disk Performance.
3. CPU Performance (Operations executed per second).

All the programs are designed in following structure -:

- The program when executed will ask to enter the number of threads and based upon the input it will run the program with given threads.
- The various parameters like latency, throughput, IOPS, FLOPS are calculated and displayed on console window at the end of execution.
- All the programs are in Java and I have used the inbuilt Java timer for calculating the time taken for execution.
- In network when the program is executed it will ask for the number of threads to run and will give the throughput and latency of varying size block data. In disk performance it will ask for entering the number of threads and will give the latency and throughput of sequential write, sequential read, random write, random read. In CPU there are three programs giving IOPS, FLOPS and giving 600 samples for running the program for 10 minutes.
- All the benchmarking was done by running programs on Amazon EC2 micro instance and for networking the benchmarking was done by connecting two different micro instance of Amazon EC2.

For performing the benchmarking of the different part of computer the program is designed as follows-:

1. Networking –:
 - The network part is divided into two parts TCP and UDP. Each UDP and TCP have server and client program which is written in java.
 - The program is designed such that it supports the multithreading and based on the user input how many threads he want to create the server opens the server socket on the different port and continuously listen for the incoming client request on that port .

- The client side program will ask for IP address and port number of the server we want to connect and will create a socket connection in TCP and in UDP it will create a datagram Packet to be send to the server.
- For sending the data the string of specified size is created at client side and send wrapping as object in TCP and datagram packet in UDP
- Both TCP and UDP client program will send the data of size 1 Byte, 1 KiloByte and 64 KiloByte to the server in TCP as in form of object and in UDP inform of packet and the server send back the same data to the client giving us the round trip time calculated at the client side.
- Throughput in Mbps and latency in milliseconds are calculated and displayed in console.
- For two thread the client will start two different ports and will send data into two threads. So we will get two different values for throughput and latency for each thread. In that case we have to take the biggest value of throughput and smallest value of latency.

2. Disk Performance :-

- The program is designed to perform the four function to measure the disk performance sequential write, sequential read, random write and random read.
- The program is written in java with multithreading support and the threads are started in the constructor based on the user input.
- Program is written in Java and it consists of four methods and they are called in the run function which is executed each time we start the new thread.
- Throughput is calculated in MegaByte per second and latency is calculated in milliseconds.
- For two threads it will give the throughput and latency of two different threads and we have to take the average of the value.

3. CPU Performance :-

- The program is designed to do the various operations like multiplication and addition supporting multithreading. The program displays the values of instruction executed per second in Giga-FLOPS and Giga-IOPS. Two different programs are written to calculate the Integer Operations Per Second(IOPS) and Floating Point Operations per second(FLOPS). The inbuilt timer in java is used for calculating the time. The duration for calculating the operation is calculated in the main thread.
- The user has to pass the number of threads to run and based on that it will create the array of the threads for executing the operations.
- A separate java program is designed for getting the number of operations performed from 4 threads and then adding it up for and getting sample for every second taking 600 samples and plotting the graph for it.

Improvement and Extension to the program-:

The further improvements and modification that we can do in our program can be

Network Performance -:

- We can increase the concurrency level up to 4 and 8 threads to compare the throughput and latency.
- We can increase the buffer size from 64KB to 1MB and 100MB and measure various performance parameters for better comparison.

Disk Performance -:

- We can increase the concurrency from 2 threads to 4 and 8 threads.
- In my program I am writing with the buffer size of 1 Byte, 1 KB and 1 MB we can increase the buffer size to 10MB and 100MB and compare the performance.
- In the disk there is possibility of having the bad sectors resulting in increased latency so we have to take in account for that too.

CPU Performance -:

- The CPU program is doing the simple arithmetic operations like addition and multiplication without much dependencies. We can increase the dependencies and make the operations more complex and calculating the GIOPS and GFLOPS values.
- The benchmarking is performed only on the ec2 micro instance so we can perform it on various other instance and compare the results to calculate the variations based on the hardware.

Performance Evaluation

. All the programs are executed for three times and the average values is taken as the results. First the environment setup of the system on which the benchmarking is performed is described. The detailed description of each part is described and finally the average data values are described in the table and then plotted in form of graph for better visualization.

System Environment -:

All the experiments are performed on Amazon EC2 t2 micro instance in Ubuntu AMI. It has 1 core with 1GB RAM and Intel Xeon processor with 2.5 GHz speed.

Overview of the Experiments -:

- Network Benchmarking is done for both TCP and UDP for varying concurrency of up to two threads. The experiments are performed using two AWS t2 micro instance and the latency and throughput are calculated for threads(1,2) for block size of 1 Byte, 1 KB and 64 KB.
- Disk Benchmarking is for four different parameters of sequential write, sequential read, random write and random read with varying concurrency of 1 and 2 threads. The average throughput and latency value is calculated for each parameter.
- CPU Benchmarking is done to calculate the giga-IOPS and giga-FLOPS for varying concurrency of 1,2 and 4 threads. The separate program is executed to get the 600 samples for every second.

Results and Analysis -:


Network -:


- The experiment is performed three times and then taken the average to get the consistent value.
- Two EC2 t2micro instance are used for both TCP and UDP. One instance the server is started on different ports and it continuously listen to the incoming request from the client. On the other instance client is started and is connected to the server.
- The throughput and latency values are calculated for various parameters for TCP and UDP which clearly states that the throughput of UDP is a way more when compared to TCP.
- The following table shows the different values for TCP and UDP for 1 and 2 threads.

Throughput Result -:

Size of blocks	Average Throughput TCP 1-Thread (Mbps)	Average Throughput UDP 1-Thread (Mbps)	Average Throughput TCP 2-Threads (Mbps)	Average Throughput UDP 2-Threads (Mbps)
1 B	0.0015	0.0554	0.0065	0.1005
1 KB	0.2044	263.2331	0.3985	400.43
64 KB	33.6037	10298.3555	37.546	20012.34

Size of blocks	Standard - Deviation Throughput TCP 1-Thread (Mbps)	Standard - Deviation Throughput UDP 1-Thread (Mbps)	Standard - Deviation Throughput TCP 2-Threads (Mbps)	Standard -Deviation Throughput UDP 2-Threads (Mbps)
1 B	0.0024	0.0562	0.0031	0.0103
1 KB	0.0193	30.356	0.1029	48.43
64 KB	3.965	207.56	5.673	107.59






Analysis -:


- As seen in the graph with increasing the packet size the throughput is increased with a larger values.
- Increasing the thread increases the throughput almost to double the value of throughput obtained from single thread.
- Comparing the UDP and TCP we can easily see that the throughput of UDP is having a much greater value as compared to TCP. TCP is connection oriented protocol compared to UDP which is connectionless protocol, so there will be a lot of overhead in TCP compared to UDP decreasing the throughput.

Latency Results -:

Size of blocks	Latency TCP 1-Thread (millisecond)	Latency UDP 1-Thread (millisecond)	Latency TCP 2-Threads (millisecond)	Latency UDP 2-Threads (millisecond)
1 B	650466.03	18051.181	658095.34	19003.56
1 KB	4892.24	3.7989	4842.01	3.571
64 KB	29.76	0.0971	47.34	0.0801

Size of blocks	Standard-Deviation Latency TCP 1-Thread (millisecond)	Standard-Deviation Latency UDP 1-Thread (millisecond)	Standard-Deviation Latency TCP 2-Threads (millisecond)	Standard-Deviation Latency UDP 2-Threads (millisecond)
1 B	90653.67	349.71	98653.56	2561.76
1 KB	308.45	1.563	342.78	1.31
64 KB	2.47	0.1045	3.71	0.134





UDP output -:

```

ec2-user@ip-172-31-31-82:~$ javac clientUDP.java
ec2-user@ip-172-31-31-82:~$ java clientUDP
Enter the number of threads you wanna start:-
1
Enter the ip address of the server:-
172.31.27.132
The throughput of size 1 is 0.0554 MBps
The latency of size 1 is 18051.1810 ms
The throughput of size 1024 is 263.2331 MBps
The latency of size 1024 is 3.7989 ms
The throughput of size 63535 is 10298.3555 MBps
The latency of size 63535 is 0.0971 ms
ec2-user@ip-172-31-31-82:~$ java clientUDP
Enter the number of threads you wanna start:-
2
Enter the ip address of the server:-
172.31.27.132
The throughput of size 1 is 0.0525 MBps
The throughput of size 1 is 0.2771 MBps
The latency of size 1 is 3608.3488 ms
The throughput of size 1024 is 288.0873 MBps
The latency of size 1024 is 3.4712 ms
The latency of size 1 is 19059.3219 ms
The throughput of size 1024 is 274.0025 MBps
The latency of size 1024 is 3.6496 ms
The throughput of size 63535 is 12137.4488 MBps
The latency of size 63535 is 0.0824 ms
The throughput of size 63535 is 12806.6928 MBps
The latency of size 63535 is 0.0781 ms
ec2-user@ip-172-31-31-82:~$
  
```


Analysis -:

- Analyzing the graph we can see that the latency remains almost same even if we increase the threads.
- Comparing TCP and UDP we can easily see that latency of TCP is having a bigger value compared to UDP because of the overhead in TCP protocol for establishing the secure connection for data transfer.
- Also as increase the data size the latency decreases at a greater value. The bigger the data size the less is the latency.

Extra Credit -:

Performing benchmarking through iperf in AWS micro instance

Server Side

install iperf package in the system

Run the server

iperf3 -s

Server listening on 5201

Accepted connection from ::1, port 56391

[5] local ::1 port 5201 connected to ::1 port 56392

[ID] Interval Transfer Bandwidth

[5] 0.00-1.00 sec 4.89 GBytes 42.0 Gbits/sec

[5] 1.00-2.00 sec 5.07 GBytes 43.5 Gbits/sec

[5] 2.00-3.00 sec 5.04 GBytes 43.3 Gbits/sec

[5] 3.00-4.00 sec 4.84 GBytes 41.6 Gbits/sec

[5] 4.00-5.00 sec 4.89 GBytes 42.0 Gbits/sec

[5] 4.00-5.00 sec 4.89 GBytes 42.0 Gbits/sec

[ID] Interval Transfer Bandwidth

[5] 0.00-5.00 sec 0.00 Bytes 0.00 bits/sec

sender

[5] 0.00-5.00 sec 29.0 GBytes 49.8 Gbits/sec

receiver

iperf3: the client has terminated

Server listening on 5201

Accepted connection from 172.31.72.5, port 56061

[5] local 172.31.72.5 port 5201 connected to 172.31.72.5 port 56062

[ID] Interval Transfer Bandwidth

[5] 0.00-1.00 sec 4.80 GBytes 41.3 Gbits/sec

[5] 1.00-2.00 sec 4.95 GBytes 42.5 Gbits/sec

[5] 2.00-3.00 sec 4.98 GBytes 42.7 Gbits/sec

[5] 3.00-4.00 sec 4.99 GBytes 42.8 Gbits/sec

[5] 4.00-5.00 sec 4.84 GBytes 41.5 Gbits/sec

[5] 5.00-6.00 sec 4.84 GBytes 41.6 Gbits/sec

[5] 6.00-7.00 sec 4.87 GBytes 41.8 Gbits/sec

[5] 7.00-8.00 sec 4.94 GBytes 42.4 Gbits/sec

[5] 8.00-9.00 sec 4.97 GBytes 42.6 Gbits/sec

```
[ 5] 9.00-10.00 sec 4.93 GBytes 42.4 Gbits/sec
[ 5] 10.00-10.04 sec 134 MBytes 29.9 Gbits/sec
-----
[ ID] Interval      Transfer   Bandwidth   Retr
[ 5] 0.00-10.04 sec 49.2 GBytes 42.1 Gbits/sec 0      sender
[ 5] 0.00-10.04 sec 49.2 GBytes 42.1 Gbits/sec      receive
```

Client Side -:

```
iperf3 -c 172.31.72.5
Connecting to host 172.31.72.5, port 5201
[ 4] local 172.31.72.5 port 56062 connected to 172.31.72.5 port 5201
[ ID] Interval      Transfer   Bandwidth   Retr Cwnd
[ 4] 0.00-1.00 sec 4.94 GBytes 42.5 Gbits/sec 0 3.18 MBytes
[ 4] 1.00-2.00 sec 4.95 GBytes 42.5 Gbits/sec 0 3.18 MBytes
[ 4] 2.00-3.00 sec 4.98 GBytes 42.7 Gbits/sec 0 3.18 MBytes
[ 4] 3.00-4.00 sec 4.99 GBytes 42.9 Gbits/sec 0 3.18 MBytes
[ 4] 4.00-5.00 sec 4.83 GBytes 41.5 Gbits/sec 0 3.18 MBytes
[ 4] 5.00-6.00 sec 4.84 GBytes 41.6 Gbits/sec 0 3.18 MBytes
[ 4] 6.00-7.00 sec 4.87 GBytes 41.8 Gbits/sec 0 3.18 MBytes
[ 4] 7.00-8.00 sec 4.94 GBytes 42.4 Gbits/sec 0 3.18 MBytes
[ 4] 8.00-9.00 sec 4.96 GBytes 42.6 Gbits/sec 0 3.18 MBytes
[ 4] 9.00-10.00 sec 4.92 GBytes 42.3 Gbits/sec 0 3.18 MBytes
-----
[ ID] Interval      Transfer   Bandwidth   Retr
[ 4] 0.00-10.00 sec 49.2 GBytes 42.3 Gbits/sec 0      sender
[ 4] 0.00-10.00 sec 49.2 GBytes 42.3 Gbits/sec      receiver
```

The maximum throughput we are getting is 42.2 Gbits/s

The efficiency when compared to the ipref is given as

Efficiency = (Calculated maximum throughput*100/maximum troughput in Ipref)

$$= (19.54*100)/42.2$$

Efficiency = 46.30

Disk Performance -:

- The experiment was performed three times and its average was taken to make the result more consistent and the graphs are plotted on that results.
- All the experiments were performed on Amazon t2 micro instance.

- The results will give throughput and latency as the output for all the four functions sequential write, sequential read, random write and random read for 1 and 2 threads.
- The data is plotted in the table and for better visualization it is plotted in graphs.

For Sequential write:

Block size	Average(Throughput-Sequential Write)	Standard Deviation(Throughput-Sequential Write)	Average(Latency-Sequential write)	Standard Deviation((Latency-Sequential Write)
1 Byte(1 thread)	42.46	21.3	23.55	10.13
1 KB (1 thread)	349.46	103.5	2.86	1.42
1 MB(1 Thread)	229.72	109.7	4.35	3.42
1 Byte(2 Thread)	36.67	14.6	27.63	18.71
1 KB (2 Thread)	439.29	102.5	2.38	1.12
1 MB(2 Thread)	367.54	131.7	3.15	2.17

For Sequential Read:

Number of Thread	Average(Throughput-Sequential Read)	Standard Deviation(Throughput-Sequential Read)	Average(Latency-Sequential Read)	Standard Deviation((Latency-Sequential Read)
1 Byte(1 thread)	3.18	2.17	314.79	117.3
1 KB (1 thread)	739.4	302.1	1.35	1.07
1 MB(1 Thread)	1357.52	438.98	0.74	0.53
1 Byte(2 Thread)	3.2	2.13	309.51	102.5
1 KB (2 Thread)	1396.22	432.54	0.89	0.74
1 MB(2 Thread)	2010.14	783.23	0.76	0.94

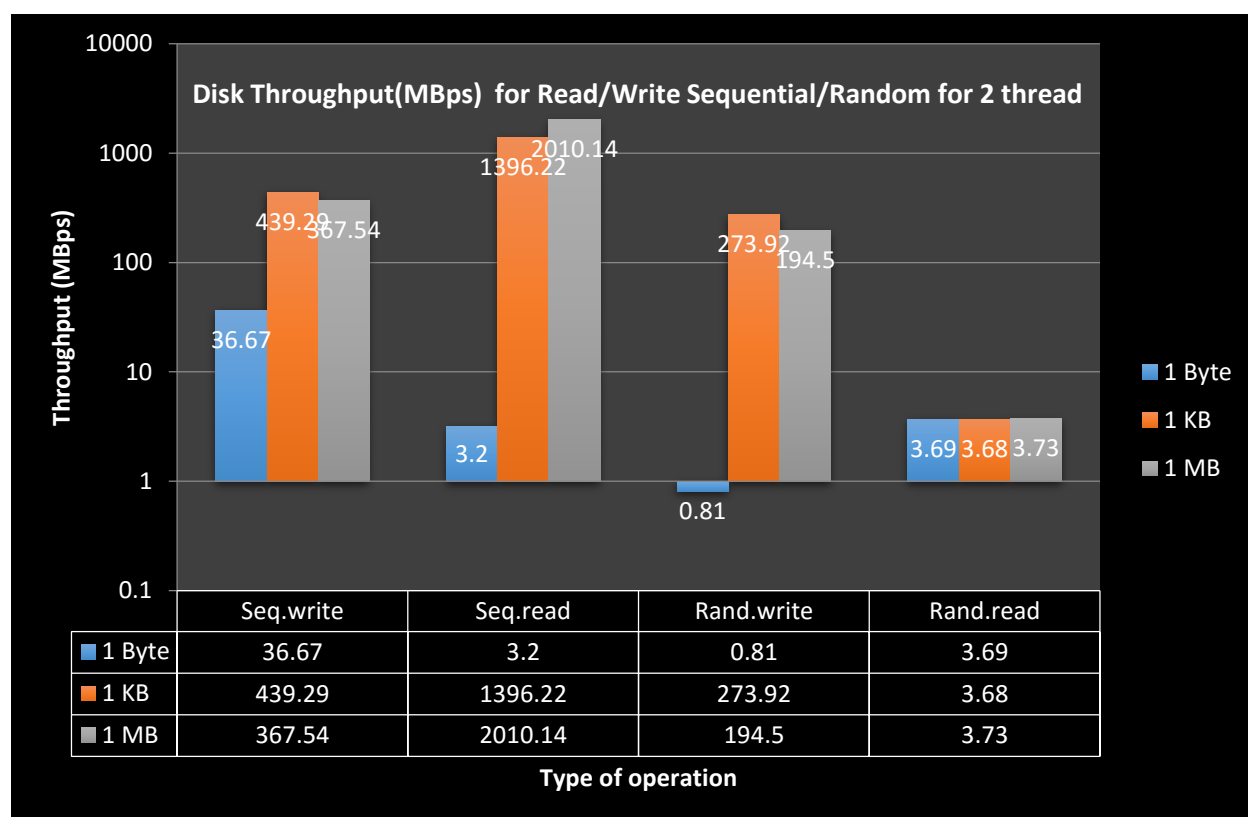

For Random write :

Number of Thread	Average(Throughput-Random write)	Standard Deviation(Throughput-Random write)	Average(Latency-Random write)	Standard Deviation((Latency-Random write)
1 Byte(1 thread)	0.79	0.93	1265.48	234.6
1 KB (1 thread)	283.54	108.42	3.53	1.76
1 MB(1 Thread)	213.81	78.56	4.68	2.35
1 Byte(2 Thread)	0.81	0.34	1238.25	308.78
1 KB (2 Thread)	273.92	103.56	3.71	1.39
1 MB (2 Thread)	256.76	106.56	5.18	2.87

For Random Read -:

Number of Thread	Average(Throughput-Random Read)	Standard Deviation(Throughput-Random Read)	Average(Latency-Random Read)	Standard Deviation((Latency-Random Read)
1 Byte(1 thread)	3.51	0.87	284.66	76.45
1 KB (1 thread)	3.56	0.93	280.9	67.87
1 MB(1 Thread)	3.7	1.03	270.37	79.67
1 Byte(2 Thread)	3.69	0.56	274.38	98.56
1 KB (2 Thread)	3.68	0.46	277.22	87.56
1 MB(2 Thread)	3.73	0.78	268.39	94.53


Throughput Results -:

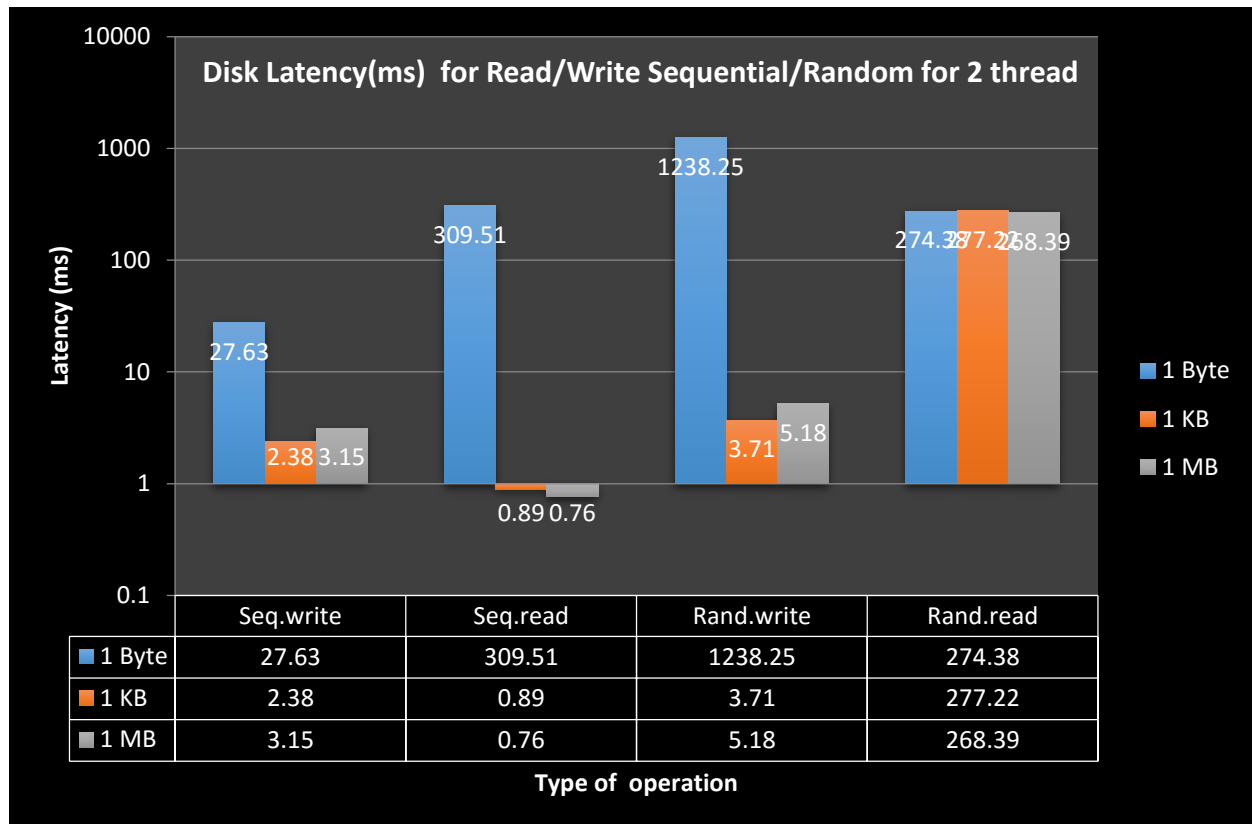


Throughput Analysis -:

- Comparing the graph and the table in the graph the throughput results for the sequential write is much greater then the random write as in sequential write the data is written sequentially whereas in random first we have to place the pointer and then write the data decreasing the throughput.
- Similarly, for reading the sequential read is having much greater value compared to the random read because for random read the data is stored in form of byte array and so we have to first jump to the position from where the data is to be read and then read the data creating a lot of overhead decreasing the throughput at a greater extent.
- For random read the throughput value is very less compared to the sequential read as it has to set the pointer every time to new position and then read creating an overhead resulting in the performance.
- For two threads the value of the sequential read and write increases whereas for random read and write it is almost the same.
- It can be seen from the graph increasing the buffer size does not increase the throughput. It will increase the throughput to a certain size and then the value will go down.

Latency Results -:





Latency Analysis -:

- Comparing the graph we can see that the values of the latency decreases as we increase the buffer size. The latency for the sequential write is lower than the random write and similarly for sequential read the value is lower than the random read.
- Increasing the thread does not decrease the latency to a larger extent but a minor change is observed.

Extra Credit -:

IOZONE

```
sudo ./iozone -a -s 1024
iozone: Performance Test of File I/O
  Version $Revision: 3.283 $
Compiled for 32 bit mode.
Build: linux
```

Contributors: William Norcott, Don Capps, Isom Crawford, Kirby Collins
 Al Slater, Scott Rhine, Mike Wisner, Ken Goss
 Steve Landherr, Brad Smith, Mark Kelly, Dr. Alain CYR,
 Randy Dunlap, Mark Montague, Dan Million,
 Jean-Marc Zucconi, Jeff Blomberg, Benny Halevy,
 Erik Habbinga, Kris Strecker, Walter Wong.

Run began: Fri Feb 12 22:07:36 2016


Efficiency = 74.93

Sample Output of Disk Program -:

```
ec2-user@ip-172-31-27-132:~
The size of the file is 10MB
Enter the number of threads you want to create-:
2
The throughput of the sequential write for buffer size 1bytes is 40.91 MBps
The latency of the sequential write for buffer size 1bytes is 24.44 ms
The throughput of the sequential write for buffer size 1024bytes is 345.43 MBps
The latency of the sequential write for buffer size 1024bytes is 2.89 ms
The throughput of the sequential write for buffer size 1048576bytes is 231.49 MBps
The latency of the sequential write for buffer size 1048576bytes is 4.32 ms
The throughput of the sequential write for buffer size 1bytes is 32.44 MBps
The latency of the sequential write for buffer size 1bytes is 30.82 ms
The throughput of the sequential write for buffer size 1024bytes is 533.15 MBps
The latency of the sequential write for buffer size 1024bytes is 1.08 ms
The throughput of the sequential write for buffer size 1048576bytes is 503.59 MBps
The latency of the sequential write for buffer size 1048576bytes is 1.99 ms
-----Sequential Read-----
The throughput of the sequential read for buffer size 1bytes is 3.15 MBps
The latency of the sequential read for buffer size 1bytes is 317.14 ms
The throughput of the sequential read for buffer size 1024bytes is 777.73 MBps
The latency of the sequential read for buffer size 1024bytes is 1.29 ms
The throughput of the sequential read for buffer size 1048576bytes is 1310.41 MBps
The latency of the sequential read for buffer size 1048576bytes is 0.76 ms
-----Sequential Read-----
The throughput of the sequential read for buffer size 1bytes is 3.31 MBps
The latency of the sequential read for buffer size 1bytes is 301.88 ms
The throughput of the sequential read for buffer size 1024bytes is 2014.71 MBps
The latency of the sequential read for buffer size 1024bytes is 0.50 ms
The throughput of the sequential read for buffer size 1048576bytes is 2709.88 MBps
The latency of the sequential read for buffer size 1048576bytes is 0.37 ms
-----Random Write-----
The throughput of the random write for buffer size 1bytes is 0.81 MBps
The latency of the random write for buffer size 1bytes is 1238.25 ms
The throughput of the random write for buffer size 1024bytes is 269.69 MBps
The latency of the random write for buffer size 1024bytes is 3.71 ms
The throughput of the random write for buffer size 1048576bytes is 193.01 MBps
The latency of the random write for buffer size 1048576bytes is 5.18 ms
-----Random Write-----
The throughput of the random write for buffer size 1bytes is 0.81 MBps
The latency of the random write for buffer size 1bytes is 1230.14 ms
The throughput of the random write for buffer size 1024bytes is 273.92 MBps
The latency of the random write for buffer size 1024bytes is 3.65 ms
The throughput of the random write for buffer size 1048576bytes is 196.83 MBps
The latency of the random write for buffer size 1048576bytes is 5.08 ms
-----Random Read-----
The throughput of the random read for buffer size 1bytes is 3.64 MBps
The latency of the random read for buffer size 1bytes is 274.38 ms
The throughput of the random read for buffer size 1024bytes is 3.61 MBps
The latency of the random read for buffer size 1024bytes is 277.22 ms
The throughput of the random read for buffer size 1048576bytes is 3.73 MBps
The latency of the random read for buffer size 1048576bytes is 268.39 ms
-----Random Read-----
The throughput of the random read for buffer size 1bytes is 3.74 MBps
The latency of the random read for buffer size 1bytes is 267.50 ms
The throughput of the random read for buffer size 1024bytes is 3.76 MBps
The latency of the random read for buffer size 1024bytes is 265.82 ms
The throughput of the random read for buffer size 1048576bytes is 3.73 MBps
The latency of the random read for buffer size 1048576bytes is 267.90 ms
[ec2-user@ip-172-31-27-132 ~]$
```

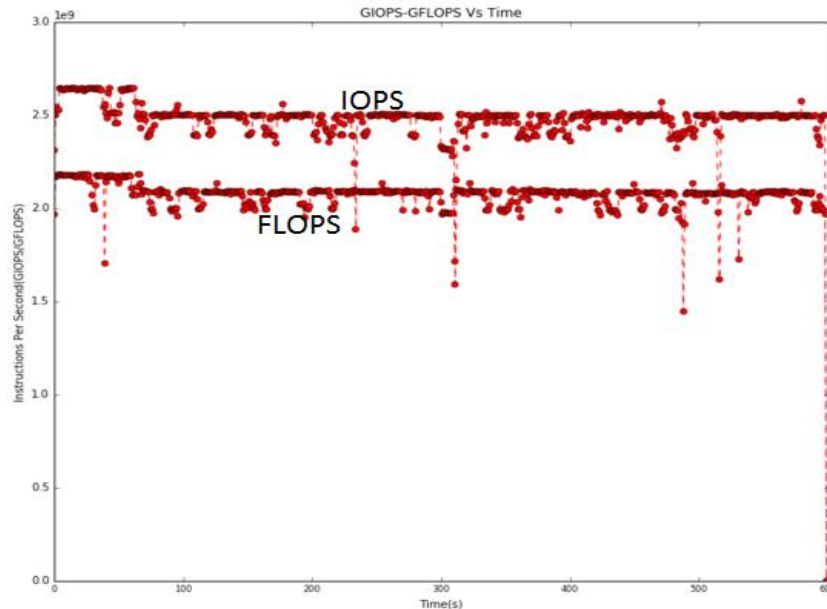
CPU Performance -:

- The experiment was performed three times and its average was taken to make the result more consistent and the graphs are plotted on that results.
- All the experiments were performed on Amazon t2 micro instance.
- The output is measured in gig-IOPS and giga-FLOPS. The performance java program will run for 10 minutes and will create a file in which all the 600 samples taken at every second is stored.



- Analyzing the graph we can see that with increase in the threads the FLOPS and IOPS are decreasing. The main reason behind this is because we are having the system of one core and also we have not taken the multiplication factor with number of threads.
- Compared to IOPS the FLOPS are having smaller value because floating point operations take more time compared to integer.
- The program is having a lot of dependencies which result in lesser number of operations executed.

The following graph is plotted by taking the value of operations performed every second.



- The graph shows the 600 samples value of the operations performed per second in GIOPS FLOPS.
- Comparing the result, we can see that the IOPS values are greater than FLOPS
- We can see that the value obtained is approximately equal to 2.5 IOPS and 2.3 FLOPS.

The theoretical performance can be calculated as follows :

$$= (\text{CPU speed in GHz}) * (\text{number of CPU cores}) * (\text{CPU instructions per node})$$

$$= 2.5 * 1 * 8 * 2$$

$$= 40 \text{ GFLOPS/sec.}$$

Efficiency -:

$$= (\text{Obtained Performance} * 100) / \text{Theoretical performance}$$

$$= 11.77 \%$$

The theoretical efficiency is the ideal efficiency when no process are running and calculating the simplest floating point operations without any dependencies. In my program I have made many dependencies which makes next to wait for previous operation and as the reason the compute cycles are wasted.

Extra Credit -:

The linpack screenshot is as follows -:

```
ec2-user@ip-172-31-50-177:~$ cd /tmp/linpack/ && ./runme_xeon4
This is a SAMPLE run script for SMP LINPACK. Change it to reflect
the correct number of CPUs/threads, problem input files, etc..
Fri Feb 12 09:50:03 UTC 2016
Intel(R) Optimized LINPACK benchmark data
Current date/time: Fri Feb 12 09:50:03 2016

CPU frequency: 2.690 GHz
Number of CPUs: 1
Number of cores: 1
Number of threads: 1

Parameters are set to:
Number of tests: 15
Number of equations to solve (problem size) : 1000 2000 5000 10000 15000 18000 20000 22000 25000 26000 27000 30000 35000 40000 45000
Leading dimension of array : 1000 2000 5000 10000 15000 18000 20016 22000 25000 26000 27000 30000 35000 40000 45000
Number of trials to run : 4 2 2 2 2 2 2 2 2 2 1 1 1 1 1
Data alignment value (in kbytes) : 4 4 4 4 4 4 4 4 4 4 1 1 1 1 1
Maximum memory requested that can be used=808204096, at the size=10000

===== Timing linear equation system solver =====
Size LDA Align. Time(s) GFlops Residual Residual(norm) Check
1000 1000 4 0.026 26.1091 9.632295e-13 3.284860e-02 pass
1000 1000 4 0.025 26.3512 9.632295e-13 3.284860e-02 pass
1000 1000 4 0.026 26.0333 9.632295e-13 3.284860e-02 pass
1000 1000 4 0.025 26.4146 9.632295e-13 3.284860e-02 pass
2000 2000 4 0.192 27.0090 4.746648e-12 4.129002e-02 pass
2000 2000 4 0.191 28.0153 4.746648e-12 4.129002e-02 pass
5000 5000 4 2.485 33.5495 2.651185e-11 3.696863e-02 pass
5000 5000 4 2.464 33.8342 2.651185e-11 3.696863e-02 pass
10000 10000 4 18.323 36.3050 9.014595e-11 3.178637e-02 pass
10000 10000 4 18.314 36.4122 9.014595e-11 3.178637e-02 pass

Performance Summary (GFlops)
Size LDA Align. Average Maximal
1000 1000 4 26.2420 26.4146
2000 2000 4 27.9124 28.0153
5000 5000 4 33.6918 33.8342
10000 10000 4 36.4041 36.4122

Residual checks PASSED

End of tests
Done: Fri Feb 12 09:50:52 UTC 2016
[ec2-user@ip-172-31-50-177:~]$
```

- The maximum value of GFLOPS obtained will be 36.4122 in amazon EC2 t2 micro instance while running the linpack.